

# NNI学生项目2020

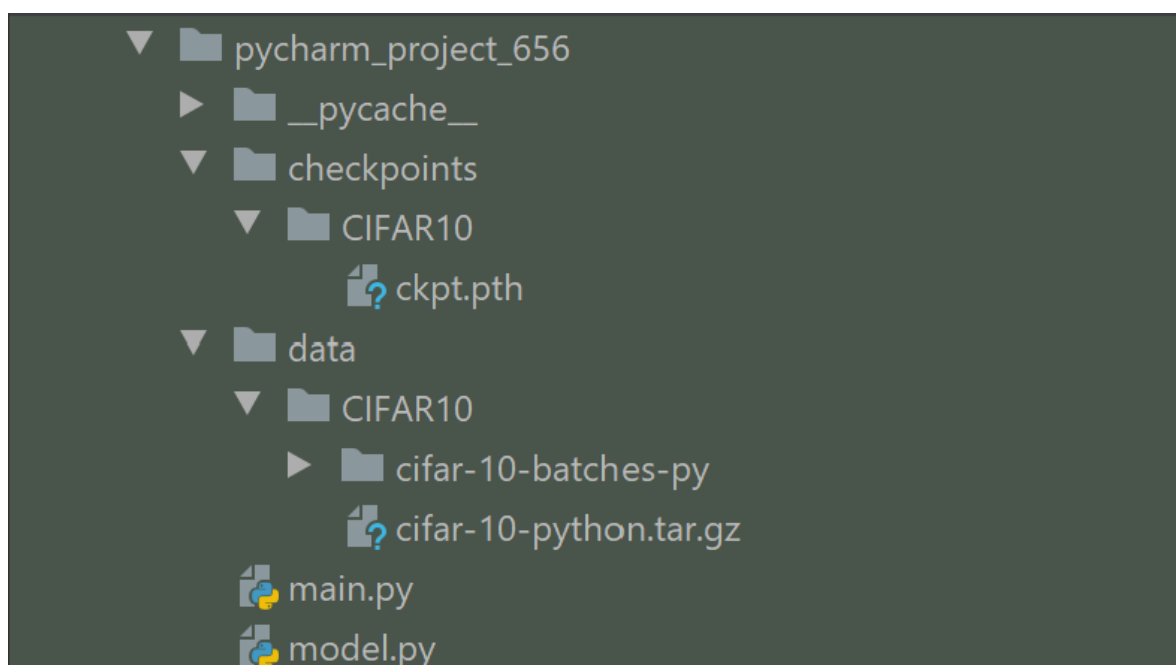
## Task 1.2.1

ID: 14 Name: 中科大飙车队

### 任务描述

- 访问Python机器学习库网站
- 完成初级任务: Training a classifier
- 以CIFAR10为样例, 通过图像分类, 训练小型神经网络, 加深同学对PyTorch Tensor Library和神经网络的理解, 帮助初级同学补充机器学习相关的背景知识。
- 提交相关文档, 包括配置文件、代码、实现、结果。

### 代码结构



- checkpoints/: 保存模型参数
- data/CIFAR10/: 数据集所在地
- main.py: 主程序的运行
- model.py: pytorch模型所在地

### 模型参数

我们使用的模型是类似于VGG架构但比VGG模型参数要少的模型(后文中称为smallVGG). VGG的基本结构是由两个卷积层+一个池化层组成的。我们使用的smallVGG由三个这样的结构组成。具体的结构我们使用torchsummary工具包进行呈现如下:

Layer (type)	Output Shape	Param #
Conv2d-1	[32, 64, 64, 64]	1,792
ReLU-2	[32, 64, 64, 64]	0
Conv2d-3	[32, 64, 64, 64]	36,928

ReLU-4	[32, 64, 64, 64]	0
MaxPool2d-5	[32, 64, 32, 32]	0
Conv2d-6	[32, 128, 32, 32]	73,856
ReLU-7	[32, 128, 32, 32]	0
Conv2d-8	[32, 128, 32, 32]	147,584
ReLU-9	[32, 128, 32, 32]	0
MaxPool2d-10	[32, 128, 16, 16]	0
Conv2d-11	[32, 128, 16, 16]	147,584
ReLU-12	[32, 128, 16, 16]	0
Conv2d-13	[32, 128, 16, 16]	147,584
ReLU-14	[32, 128, 16, 16]	0
MaxPool2d-15	[32, 128, 8, 8]	0
Linear-16	[32, 512]	1,049,088
Sigmoid-17	[32, 512]	0
Linear-18	[32, 10]	5,130

=====  
Total params: 1,609,546

Trainable params: 1,609,546

Non-trainable params: 0

-----  
Input size (MB): 1.50

Forward/backward pass size (MB): 442.25

Params size (MB): 6.14

Estimated Total Size (MB): 449.89  
-----

我们选取batch size为32.选取的优化器为Adam优化器，学习率为0.001而其他参数符合默认。我们每20epochs对学习率进行减半。

## 代码实现

model.py

```
#import *** omitted for simplicity
class smallVGG(nn.Module):
    def __init__(self, NChannels):
        super(smallVGG, self).__init__()
        self.features = []
        self.layerDict = collections.OrderedDict()

        self.conv11 = nn.Conv2d(
            in_channels = NChannels,
            out_channels = 64,
            kernel_size = 3,
            padding = 1
        )
        self.features.append(self.conv11)
        self.layerDict['conv11'] = self.conv11

        self.ReLU11 = nn.ReLU()
        self.features.append(self.ReLU11)
        self.layerDict['ReLU11'] = self.ReLU11

        self.conv12 = nn.Conv2d(
            in_channels = 64,
            out_channels = 64,
            kernel_size = 3,
```

```
        padding = 1
    )
    self.features.append(self.conv12)
    self.layerDict['conv12'] = self.conv12

    self.ReLU12 = nn.ReLU()
    self.features.append(self.ReLU12)
    self.layerDict['ReLU12'] = self.ReLU12

    self.pool1 = nn.MaxPool2d(2,2)
    self.features.append(self.pool1)
    self.layerDict['pool1'] = self.pool1

    self.conv21 = nn.Conv2d(
        in_channels = 64,
        out_channels = 128,
        kernel_size = 3,
        padding = 1
    )
    self.features.append(self.conv21)
    self.layerDict['conv21'] = self.conv21

    self.ReLU21 = nn.ReLU()
    self.features.append(self.ReLU21)
    self.layerDict['ReLU21'] = self.ReLU21

    self.conv22 = nn.Conv2d(
        in_channels = 128,
        out_channels = 128,
        kernel_size = 3,
        padding = 1
    )
    self.features.append(self.conv22)
    self.layerDict['conv22'] = self.conv22

    self.ReLU22 = nn.ReLU()
    self.features.append(self.ReLU22)
    self.layerDict['ReLU22'] = self.ReLU22

    self.pool2 = nn.MaxPool2d(2,2)
    self.features.append(self.pool2)
    self.layerDict['pool2'] = self.pool2

    self.conv31 = nn.Conv2d(
        in_channels = 128,
        out_channels = 128,
        kernel_size = 3,
        padding = 1
    )
    self.features.append(self.conv31)
    self.layerDict['conv31'] = self.conv31

    self.ReLU31 = nn.ReLU()
    self.features.append(self.ReLU31)
    self.layerDict['ReLU31'] = self.ReLU31

    self.conv32 = nn.Conv2d(
```

```

        in_channels = 128,
        out_channels = 128,
        kernel_size = 3,
        padding = 1
    )
    self.features.append(self.conv32)
    self.layerDict['conv32'] = self.conv32

    self.ReLU32 = nn.ReLU()
    self.features.append(self.ReLU32)
    self.layerDict['ReLU32'] = self.ReLU32

    self.pool3 = nn.MaxPool2d(2,2)
    self.features.append(self.pool3)
    self.layerDict['pool3'] = self.pool3

    self.classifier = []

    self.feature_dims = 4 * 4 * 128
    self.fc1 = nn.Linear(self.feature_dims, 512)
    self.classifier.append(self.fc1)
    self.layerDict['fc1'] = self.fc1

    self.fc1act = nn.Sigmoid()
    self.classifier.append(self.fc1act)
    self.layerDict['fc1act'] = self.fc1act

    self.fc2 = nn.Linear(512, 10)
    self.classifier.append(self.fc2)
    self.layerDict['fc2'] = self.fc2

def forward(self, x):
    for layer in self.features:
        x = layer(x)

    x = x.view(-1, self.feature_dims)

    for layer in self.classifier:
        x = layer(x)
    return x

```

代码的结构比较简单。在定义网络(尤其是针对图像类任务)时一般为了使用清晰和方便都会将结构分为 features 和 classifier 两部分。features 一般包括卷积、池化、正则化等非全连接层，classifier 一般包括两到三个全连接层。这样网络结构比较清晰，也准确的描述了这些网络的功能。

main.py

```

# import *** omitted for simplicity
def setLearningRate(optimizer, lr):
    """Sets the learning rate to the initial LR decayed by 10 every 30 epochs"""
    for param_group in optimizer.param_groups:
        param_group['lr'] = lr
def train(DATASET = 'CIFAR10', NEpochs = 200,
        BatchSize = 32, learningRate = 1e-3, NDecreaseLR = 20, eps = 1e-3,
        AMSGrad = True, model_dir = "checkpoints/CIFAR10/", model_name =
"ckpt.pth", gpu = True):

```

```

print("DATASET: ", DATASET)

mu = torch.tensor([0.485, 0.456, 0.406], dtype=torch.float32)
sigma = torch.tensor([0.229, 0.224, 0.225], dtype=torch.float32)
Normalize = transforms.Normalize(mu.tolist(), sigma.tolist())
Unnormalize = transforms.Normalize((-mu / sigma).tolist(), (1.0 /
sigma).tolist())
tsf = {
    'train': transforms.Compose(
        [
            transforms.RandomHorizontalFlip(),
            transforms.RandomAffine(degrees = 10, translate = [0.1, 0.1], scale =
[0.9, 1.1]),
            transforms.ToTensor(),
            Normalize
        ]),
    'test': transforms.Compose(
        [
            transforms.ToTensor(),
            Normalize
        ])
}
trainset = torchvision.datasets.CIFAR10(root='./data/CIFAR10', train = True,
download=True, transform = tsf['train'])
testset = torchvision.datasets.CIFAR10(root='./data/CIFAR10', train = False,
download=True, transform = tsf['test'])

net = smallVGG(3)

x_train, y_train = trainset.data, trainset.targets,
x_test, y_test = testset.data, testset.targets,

trainloader = torch.utils.data.DataLoader(trainset, batch_size = BatchSize,
shuffle = True, num_workers = 1)
testloader = torch.utils.data.DataLoader(testset, batch_size = 1000,
shuffle = False, num_workers = 1)

trainIter = iter(trainloader)
testIter = iter(testloader)
criterion = nn.CrossEntropyLoss()
softmax = nn.Softmax(dim=1)

if gpu:
    net.cuda()
    criterion.cuda()
    softmax.cuda()

optimizer = optim.Adam(params = net.parameters(), lr = learningRate, eps =
eps, amsgrad = AMSGrad)

NBatch = int(len(trainset) / BatchSize)
cudnn.benchmark = True
for epoch in range(NEpochs):
    lossTrain = 0.0
    accTrain = 0.0
    for i in range(NBatch):

```

```

try:
    batchX, batchY = trainIter.next()
except StopIteration:
    trainIter = iter(trainloader)
    batchX, batchY = trainIter.next()

if gpu:
    batchX = batchX.cuda()
    batchY = batchY.cuda()

optimizer.zero_grad()
logits = net.forward(batchX)
prob = softmax(logits)

loss = criterion(logits, batchY)
loss.backward()
optimizer.step()

lossTrain += loss.cpu().detach().numpy() / NBatch
if gpu:
    pred = np.argmax(prob.cpu().detach().numpy(), axis = 1)
    groundTruth = batchY.cpu().detach().numpy()
else:
    pred = np.argmax(prob.detach().numpy(), axis = 1)
    groundTruth = batchY.detach().numpy()

acc = np.mean(pred == groundTruth)
accTrain += acc / NBatch

if (epoch + 1) % NDecreaseLR == 0:
    learningRate = learningRate / 2.0
    setLearningRate(optimizer, learningRate)
accTest, lossTest = test(testloader, net, gpu=gpu)
print("Epoch: ", epoch, "Train Loss: ", lossTrain, "Train accuracy: ",
accTrain*100, "Test Loss: ", \
    lossTest, "Test accuracy: ", accTest*100)

if not os.path.exists(model_dir):
    os.makedirs(model_dir)
torch.save(net, model_dir + model_name)
print("Model saved")
def test(testloader, net, gpu = True):
    testIter = iter(testloader)
    acc = 0.0
    NBatch = 10
    cri = nn.CrossEntropyLoss()
    lossTest = 0.0
    for i, data in enumerate(testIter, 0):
        batchX, batchY = data
        if gpu:
            batchX = batchX.cuda()
            batchY = batchY.cuda()
        logits = net.forward(batchX)
        loss = cri(logits, batchY)
        lossTest += loss.cpu().detach().numpy() / NBatch
        if gpu:
            pred = np.argmax(logits.cpu().detach().numpy(), axis = 1)
            groundTruth = batchY.cpu().detach().numpy()

```

```

else:
    pred = np.argmax(logits.detach().numpy(), axis = 1)
    groundTruth = batchY.detach().numpy()
    acc += np.mean(pred == groundTruth)
accTest = acc / NBatch
return accTest, lossTest
if __name__ == '__main__':
    import argparse
    import sys
    import traceback

    try:
        parser = argparse.ArgumentParser()
        parser.add_argument('--dataset', type = str, default = 'CIFAR10')
        parser.add_argument('--epochs', type = int, default = 100)
        parser.add_argument('--eps', type = float, default = 1e-3)
        parser.add_argument('--AMSGrad', type = bool, default = True)
        parser.add_argument('--batch_size', type = int, default = 32)
        parser.add_argument('--learning_rate', type = float, default = 1e-3)
        parser.add_argument('--decrease_LR', type = int, default = 20)

        parser.add_argument('--nogpu', dest='gpu', action='store_false')
        parser.set_defaults(gpu=True)
        args = parser.parse_args()

        model_dir = "checkpoints/" + args.dataset + '/'
        model_name = "ckpt.pth"
        train(DATASET = args.dataset, NEPOCHS = args.epochs,
              BatchSize = args.batch_size, learningRate = args.learning_rate,
              NDecreaseLR = args.decrease_LR, eps = args.eps,
              AMSGrad = args.AMSGrad, model_dir = model_dir, model_name = model_name,
              gpu = args.gpu)

    except:
        traceback.print_exc(file=sys.stdout)
        sys.exit(1)

```

这里也只是一个简单的模型的pipeline,没有涉及到复杂的操作。类似的简单模型都可以用这样的pipeline进行训练,只要定义好网络的接口即可。

## 结果展示

经过100epoches的训练,可以得到如下的结果

```

Epoch: 0 Train Loss: 1.7037 Train accuracy: 36.69 Test Loss: 1.2625 Test
accuracy: 54.31
Epoch: 1 Train Loss: 1.2118 Train accuracy: 55.91 Test Loss: 1.0374 Test
accuracy: 63.25
Epoch: 2 Train Loss: 0.9448 Train accuracy: 66.4 Test Loss: 0.7656 Test
accuracy: 73.14
Epoch: 3 Train Loss: 0.8016 Train accuracy: 71.78 Test Loss: 0.6717 Test
accuracy: 76.36
Epoch: 4 Train Loss: 0.7119 Train accuracy: 74.97 Test Loss: 0.6023 Test
accuracy: 79.53
Epoch: 5 Train Loss: 0.6424 Train accuracy: 77.54 Test Loss: 0.6355 Test
accuracy: 78.17

```

Epoch: 6 Train Loss: 0.5985 Train accuracy: 79.16 Test Loss: 0.5348 Test accuracy: 81.79

Epoch: 7 Train Loss: 0.5644 Train accuracy: 80.49 Test Loss: 0.5145 Test accuracy: 82.15

Epoch: 8 Train Loss: 0.5323 Train accuracy: 81.5 Test Loss: 0.516 Test accuracy: 82.45

Epoch: 9 Train Loss: 0.5065 Train accuracy: 82.26 Test Loss: 0.4728 Test accuracy: 83.59

Epoch: 10 Train Loss: 0.4923 Train accuracy: 82.89 Test Loss: 0.4893 Test accuracy: 83.16

Epoch: 11 Train Loss: 0.479 Train accuracy: 83.33 Test Loss: 0.5008 Test accuracy: 82.72

Epoch: 12 Train Loss: 0.4766 Train accuracy: 83.57 Test Loss: 0.4649 Test accuracy: 84.02

Epoch: 13 Train Loss: 0.4672 Train accuracy: 84.0 Test Loss: 0.5326 Test accuracy: 81.76

Epoch: 14 Train Loss: 0.4723 Train accuracy: 83.68 Test Loss: 0.4914 Test accuracy: 83.52

Epoch: 15 Train Loss: 0.4803 Train accuracy: 83.37 Test Loss: 0.4824 Test accuracy: 83.49

Epoch: 16 Train Loss: 0.4894 Train accuracy: 83.17 Test Loss: 0.4773 Test accuracy: 83.65

Epoch: 17 Train Loss: 0.4991 Train accuracy: 82.66 Test Loss: 0.5199 Test accuracy: 82.01

Epoch: 18 Train Loss: 0.5318 Train accuracy: 81.84 Test Loss: 0.5481 Test accuracy: 81.46

Epoch: 19 Train Loss: 0.5613 Train accuracy: 80.59 Test Loss: 0.5441 Test accuracy: 81.12

Epoch: 20 Train Loss: 0.4639 Train accuracy: 84.05 Test Loss: 0.4429 Test accuracy: 84.99

Epoch: 21 Train Loss: 0.4343 Train accuracy: 85.02 Test Loss: 0.4456 Test accuracy: 84.96

Epoch: 22 Train Loss: 0.4257 Train accuracy: 85.4 Test Loss: 0.4558 Test accuracy: 84.4

Epoch: 23 Train Loss: 0.4136 Train accuracy: 85.66 Test Loss: 0.447 Test accuracy: 85.19

Epoch: 24 Train Loss: 0.4029 Train accuracy: 86.24 Test Loss: 0.4494 Test accuracy: 85.16

Epoch: 25 Train Loss: 0.4091 Train accuracy: 85.92 Test Loss: 0.4318 Test accuracy: 85.16

Epoch: 26 Train Loss: 0.3998 Train accuracy: 86.02 Test Loss: 0.4106 Test accuracy: 86.35

Epoch: 27 Train Loss: 0.3897 Train accuracy: 86.51 Test Loss: 0.4222 Test accuracy: 85.76

Epoch: 28 Train Loss: 0.3928 Train accuracy: 86.29 Test Loss: 0.4308 Test accuracy: 85.47

Epoch: 29 Train Loss: 0.3863 Train accuracy: 86.6 Test Loss: 0.4446 Test accuracy: 84.9

Epoch: 30 Train Loss: 0.3822 Train accuracy: 86.87 Test Loss: 0.4364 Test accuracy: 85.53

Epoch: 31 Train Loss: 0.3809 Train accuracy: 86.81 Test Loss: 0.4173 Test accuracy: 85.95

Epoch: 32 Train Loss: 0.3899 Train accuracy: 86.5 Test Loss: 0.4429 Test accuracy: 85.24

Epoch: 33 Train Loss: 0.3811 Train accuracy: 86.84 Test Loss: 0.4209 Test accuracy: 85.9

Epoch: 34 Train Loss: 0.3792 Train accuracy: 86.88 Test Loss: 0.4112 Test accuracy: 86.29



Epoch: 35	Train Loss: 0.3732	Train accuracy: 86.97	Test Loss: 0.4239	Test accuracy: 85.81
Epoch: 36	Train Loss: 0.3745	Train accuracy: 87.06	Test Loss: 0.4159	Test accuracy: 85.59
Epoch: 37	Train Loss: 0.3732	Train accuracy: 87.16	Test Loss: 0.446	Test accuracy: 85.38
Epoch: 38	Train Loss: 0.3691	Train accuracy: 87.52	Test Loss: 0.4075	Test accuracy: 86.34
Epoch: 39	Train Loss: 0.3676	Train accuracy: 87.17	Test Loss: 0.4343	Test accuracy: 85.61
Epoch: 40	Train Loss: 0.3368	Train accuracy: 88.17	Test Loss: 0.3882	Test accuracy: 87.01
Epoch: 41	Train Loss: 0.3208	Train accuracy: 88.99	Test Loss: 0.3807	Test accuracy: 87.28
Epoch: 42	Train Loss: 0.3117	Train accuracy: 89.25	Test Loss: 0.38	Test accuracy: 87.52
Epoch: 43	Train Loss: 0.3097	Train accuracy: 89.28	Test Loss: 0.3864	Test accuracy: 87.08
Epoch: 44	Train Loss: 0.3113	Train accuracy: 89.13	Test Loss: 0.3729	Test accuracy: 87.79
Epoch: 45	Train Loss: 0.3031	Train accuracy: 89.43	Test Loss: 0.3773	Test accuracy: 87.49
Epoch: 46	Train Loss: 0.2984	Train accuracy: 89.67	Test Loss: 0.3685	Test accuracy: 87.78
Epoch: 47	Train Loss: 0.2944	Train accuracy: 89.75	Test Loss: 0.3767	Test accuracy: 87.74
Epoch: 48	Train Loss: 0.2945	Train accuracy: 89.78	Test Loss: 0.3803	Test accuracy: 87.43
Epoch: 49	Train Loss: 0.2893	Train accuracy: 89.86	Test Loss: 0.3732	Test accuracy: 87.29
Epoch: 50	Train Loss: 0.2882	Train accuracy: 89.98	Test Loss: 0.3679	Test accuracy: 87.6
Epoch: 51	Train Loss: 0.2894	Train accuracy: 89.9	Test Loss: 0.3749	Test accuracy: 87.52
Epoch: 52	Train Loss: 0.2903	Train accuracy: 89.87	Test Loss: 0.3691	Test accuracy: 87.58
Epoch: 53	Train Loss: 0.2814	Train accuracy: 90.28	Test Loss: 0.3712	Test accuracy: 87.93
Epoch: 54	Train Loss: 0.2797	Train accuracy: 90.25	Test Loss: 0.3594	Test accuracy: 88.29
Epoch: 55	Train Loss: 0.286	Train accuracy: 90.08	Test Loss: 0.3728	Test accuracy: 87.62
Epoch: 56	Train Loss: 0.2822	Train accuracy: 90.18	Test Loss: 0.3802	Test accuracy: 87.53
Epoch: 57	Train Loss: 0.2804	Train accuracy: 90.21	Test Loss: 0.3805	Test accuracy: 87.51
Epoch: 58	Train Loss: 0.2845	Train accuracy: 90.12	Test Loss: 0.3648	Test accuracy: 87.86
Epoch: 59	Train Loss: 0.2829	Train accuracy: 90.14	Test Loss: 0.3912	Test accuracy: 87.04
Epoch: 60	Train Loss: 0.2579	Train accuracy: 91.12	Test Loss: 0.3535	Test accuracy: 88.33
Epoch: 61	Train Loss: 0.2546	Train accuracy: 91.16	Test Loss: 0.3488	Test accuracy: 88.65
Epoch: 62	Train Loss: 0.247	Train accuracy: 91.42	Test Loss: 0.3557	Test accuracy: 88.45
Epoch: 63	Train Loss: 0.2487	Train accuracy: 91.21	Test Loss: 0.352	Test accuracy: 88.42

Epoch: 64	Train Loss: 0.2469	Train accuracy: 91.42	Test Loss: 0.3521	Test accuracy: 88.62
Epoch: 65	Train Loss: 0.2453	Train accuracy: 91.49	Test Loss: 0.3528	Test accuracy: 88.59
Epoch: 66	Train Loss: 0.2448	Train accuracy: 91.34	Test Loss: 0.3611	Test accuracy: 88.27
Epoch: 67	Train Loss: 0.247	Train accuracy: 91.32	Test Loss: 0.3561	Test accuracy: 88.38
Epoch: 68	Train Loss: 0.2457	Train accuracy: 91.43	Test Loss: 0.3546	Test accuracy: 88.39
Epoch: 69	Train Loss: 0.2385	Train accuracy: 91.58	Test Loss: 0.3573	Test accuracy: 88.26
Epoch: 70	Train Loss: 0.2389	Train accuracy: 91.75	Test Loss: 0.3538	Test accuracy: 88.41
Epoch: 71	Train Loss: 0.2389	Train accuracy: 91.66	Test Loss: 0.3575	Test accuracy: 88.33
Epoch: 72	Train Loss: 0.2335	Train accuracy: 91.89	Test Loss: 0.3627	Test accuracy: 88.31
Epoch: 73	Train Loss: 0.2353	Train accuracy: 91.75	Test Loss: 0.3606	Test accuracy: 88.36
Epoch: 74	Train Loss: 0.2304	Train accuracy: 91.98	Test Loss: 0.3528	Test accuracy: 88.67
Epoch: 75	Train Loss: 0.234	Train accuracy: 91.86	Test Loss: 0.3533	Test accuracy: 88.42
Epoch: 76	Train Loss: 0.2317	Train accuracy: 91.82	Test Loss: 0.3599	Test accuracy: 88.24
Epoch: 77	Train Loss: 0.2314	Train accuracy: 91.93	Test Loss: 0.3532	Test accuracy: 88.66
Epoch: 78	Train Loss: 0.2315	Train accuracy: 91.86	Test Loss: 0.3571	Test accuracy: 88.65
Epoch: 79	Train Loss: 0.2268	Train accuracy: 92.17	Test Loss: 0.3482	Test accuracy: 88.89
Epoch: 80	Train Loss: 0.2204	Train accuracy: 92.32	Test Loss: 0.349	Test accuracy: 88.78
Epoch: 81	Train Loss: 0.219	Train accuracy: 92.4	Test Loss: 0.3467	Test accuracy: 88.97
Epoch: 82	Train Loss: 0.2207	Train accuracy: 92.31	Test Loss: 0.3464	Test accuracy: 89.02
Epoch: 83	Train Loss: 0.2178	Train accuracy: 92.48	Test Loss: 0.3458	Test accuracy: 88.94
Epoch: 84	Train Loss: 0.2162	Train accuracy: 92.38	Test Loss: 0.3443	Test accuracy: 88.88
Epoch: 85	Train Loss: 0.216	Train accuracy: 92.42	Test Loss: 0.3442	Test accuracy: 88.97
Epoch: 86	Train Loss: 0.2137	Train accuracy: 92.56	Test Loss: 0.3485	Test accuracy: 88.85
Epoch: 87	Train Loss: 0.2156	Train accuracy: 92.46	Test Loss: 0.3473	Test accuracy: 88.9
Epoch: 88	Train Loss: 0.2152	Train accuracy: 92.52	Test Loss: 0.3468	Test accuracy: 88.92
Epoch: 89	Train Loss: 0.2132	Train accuracy: 92.62	Test Loss: 0.3479	Test accuracy: 88.93
Epoch: 90	Train Loss: 0.2097	Train accuracy: 92.64	Test Loss: 0.3512	Test accuracy: 88.66
Epoch: 91	Train Loss: 0.2138	Train accuracy: 92.46	Test Loss: 0.3498	Test accuracy: 88.71
Epoch: 92	Train Loss: 0.2111	Train accuracy: 92.69	Test Loss: 0.3437	Test accuracy: 88.79

```
Epoch: 93 Train Loss: 0.212 Train accuracy: 92.5 Test Loss: 0.3494 Test accuracy: 88.54
Epoch: 94 Train Loss: 0.2153 Train accuracy: 92.45 Test Loss: 0.3438 Test accuracy: 89.06
Epoch: 95 Train Loss: 0.2077 Train accuracy: 92.76 Test Loss: 0.349 Test accuracy: 88.88
Epoch: 96 Train Loss: 0.2065 Train accuracy: 92.82 Test Loss: 0.3456 Test accuracy: 88.82
Epoch: 97 Train Loss: 0.2068 Train accuracy: 92.64 Test Loss: 0.3443 Test accuracy: 88.96
Epoch: 98 Train Loss: 0.2063 Train accuracy: 92.73 Test Loss: 0.3466 Test accuracy: 88.83
Epoch: 99 Train Loss: 0.2035 Train accuracy: 92.9 Test Loss: 0.3447 Test accuracy: 89.03
```

注意到我们的训练并没有引起严重的过拟合。整体来generalization-gap处于可接受的范围。结果使用pycharm的自动保存输出在output.pdf中。

## 结果分析与结论

- 模型比较小，结果也比较合理，没有很多复杂的改进和设计
- 后面希望用NNI调优可以取得更好的效果
- 优化器的选取要进行慎重考虑，尤其注意使用weight decay, weight decay可以有效避免训练结果的"抖动"
- 注意要避免过拟合，这里使用了数据增强的方法和weight decay，之前做别的项目的时候，试过对CIFAR10不使用数据增强，结果造成严重的过拟合。对于CIFAR10这种较小的数据集尤其要注意。